

index

April 25, 2022

1 Managing Time Series Data - Lab

1.1 Introduction

In the previous lesson, you learned that time series data are everywhere and working with time series data is an important skill for data scientists!

In this lab, you'll practice your previously learned techniques to import, clean, and manipulate time series data.

The lab will cover how to perform time series analysis while working with large datasets. The dataset can be memory intensive so your computer will need at least 2GB of memory to perform some of the calculations.

1.2 Objectives

You will be able to:

- Load time series data using Pandas and perform time series indexing
- Perform data cleaning operation on time series data
- Change the granularity of a time series

1.3 Let's get started!

Import the following libraries:

- `pandas`, using the alias `pd`
- `pandas.tseries`
- `matplotlib.pyplot`, using the alias `plt`
- `statsmodels.api`, using the alias `sm`

```
[12]: # Load required libraries
import pandas as pd
import pandas.tseries
import matplotlib.pyplot as plt
import statsmodels.api as sm
```

1.4 Loading time series data

The `statsModels` library comes bundled with built-in datasets for experimentation and practice. A detailed description of these datasets can be found [here](#). Using `statsModels`, the time series

datasets can be loaded straight into memory.

In this lab, we'll use the **Atmospheric CO2 from Continuous Air Samples at Mauna Loa Observatory, Hawaii, U.S.A.**, containing CO2 samples from March 1958 to December 2001. Further details on this dataset are available [here](#).

In the following cell:

- We loaded the co2 dataset using the `.load()` method
- Converted this into a pandas DataFrame
- Renamed the columns
- Set the 'date' column as index

```
[3]: # Load the 'co2' dataset from sm.datasets
data_set = sm.datasets.co2.load()

# load in the data_set into pandas dataframe
C02 = pd.DataFrame(data=data_set['data'])
C02.rename(columns={'index': 'date'}, inplace=True)

# set index to date column
C02.set_index('date', inplace=True)

C02.head()
```

```
[3]:          co2
date
1958-03-29  316.1
1958-04-05  317.3
1958-04-12  317.6
1958-04-19  317.5
1958-04-26  316.4
```

Let's check the data type of C02 and also display the first 15 entries of C02 as our first exploratory step.

```
[9]: # Print the data type of C02
print(C02.dtypes)

# Display the first 15 rows of C02
C02.head(15)
```

```
co2    float64
dtype: object
```

```
[9]:          co2
date
1958-03-29  316.1
1958-04-05  317.3
1958-04-12  317.6
```

| | |
|------------|-------|
| 1958-04-19 | 317.5 |
| 1958-04-26 | 316.4 |
| 1958-05-03 | 316.9 |
| 1958-05-10 | NaN |
| 1958-05-17 | 317.5 |
| 1958-05-24 | 317.9 |
| 1958-05-31 | NaN |
| 1958-06-07 | NaN |
| 1958-06-14 | NaN |
| 1958-06-21 | NaN |
| 1958-06-28 | NaN |
| 1958-07-05 | 315.8 |

With all the required packages imported and the CO2 dataset as a dataframe ready to go, we can move on to indexing our data.

1.5 Date Indexing

While working with time series data in Python, having dates (or datetimes) in the index can be very helpful, especially if they are of `DatetimeIndex` type. Further details can be found [here](#).

Display the `.index` attribute of the CO2 DataFrame:

```
[13]: # Confirm that date values are used for indexing purpose in the CO2 dataset
print(CO2.index)
```

```
DatetimeIndex(['1958-03-29', '1958-04-05', '1958-04-12', '1958-04-19',
               '1958-04-26', '1958-05-03', '1958-05-10', '1958-05-17',
               '1958-05-24', '1958-05-31',
               ...,
               '2001-10-27', '2001-11-03', '2001-11-10', '2001-11-17',
               '2001-11-24', '2001-12-01', '2001-12-08', '2001-12-15',
               '2001-12-22', '2001-12-29'],
              dtype='datetime64[ns]', name='date', length=2284, freq=None)
```

The output above shows that our dataset clearly fulfills the indexing requirements. Look at the last line:

```
dtype='datetime64[ns]', length=2284, freq='W-SAT'
```

- `dtype=datetime[ns]` field confirms that the index is made of timestamp objects.
- `length=2284` shows the total number of entries in our time series data.

1.6 Resampling

Remember that depending on the nature of analytical question, the resolution of timestamps can also be changed to other frequencies. For this dataset we can resample to monthly CO2 consumption values. This can be done by using the `.resample()` method as seen in the earlier lesson.

- Group the data into buckets representing 1 month using `.resample()` method
- Call the `.mean()` method on each group (i.e. get monthly average)

- Combine the result as one row per monthly group

```
[21]: # Group the time series into monthly buckets
CO2_monthly = CO2.resample("MS")

# CO2_monthly = CO2["co2"].resample("MS")

# Take the mean of each group
CO2_monthly_mean = CO2_monthly.mean()

# Display the first 10 elements of resulting time series
CO2_monthly_mean.head(10)
```

```
[21]:
```

| | co2 |
|------------|------------|
| date | |
| 1958-03-01 | 316.100000 |
| 1958-04-01 | 317.200000 |
| 1958-05-01 | 317.433333 |
| 1958-06-01 | NaN |
| 1958-07-01 | 315.625000 |
| 1958-08-01 | 314.950000 |
| 1958-09-01 | 313.500000 |
| 1958-10-01 | NaN |
| 1958-11-01 | 313.425000 |
| 1958-12-01 | 314.700000 |

Looking at the index values, we can see that our time series now carries aggregated data on monthly terms, shown as Freq: MS.

1.6.1 Time-series Index Slicing for Data Selection

Slice our dataset to only retrieve data points that come after the year 1990.

```
[24]: # Slice the timeseries to contain data after year 1990

CO2_monthly_mean['1990':]
```

```
[24]:
```

| | co2 |
|------------|---------|
| date | |
| 1990-01-01 | 353.650 |
| 1990-02-01 | 354.650 |
| 1990-03-01 | 355.480 |
| 1990-04-01 | 356.175 |
| 1990-05-01 | 357.075 |
| ... | ... |
| 2001-08-01 | 369.425 |
| 2001-09-01 | 367.880 |
| 2001-10-01 | 368.050 |

```
2001-11-01    369.375
2001-12-01    371.020
```

```
[144 rows x 1 columns]
```

Retrieve data starting from Jan 1990 to Jan 1991:

```
[28]: # Retrieve the data between 1st Jan 1990 to 1st Jan 1991
      CO2_monthly_mean['1990-01':'1991-01-01']
```

```
[28]:          co2
      date
1990-01-01    353.650
1990-02-01    354.650
1990-03-01    355.480
1990-04-01    356.175
1990-05-01    357.075
1990-06-01    356.080
1990-07-01    354.675
1990-08-01    352.900
1990-09-01    350.940
1990-10-01    351.225
1990-11-01    352.700
1990-12-01    354.140
1991-01-01    354.675
```

1.7 Missing Values

Find the total number of missing values in the dataset.

```
[ ]: # Find the total number of missing values in the time series
```

Remember that missing values can be filled in a multitude of ways.

- Replace the missing values in `CO2_monthly_mean` with a previous valid value
- Next, check if your attempt was successful by checking for number of missing values again

```
[ ]: # Perform backward filling of missing values
      CO2_final = None

      # Find the total number of missing values in the time series
```

Great! Now your time series data are ready for visualization and further analysis.

1.8 Summary

In this introductory lab, you learned how to load and manipulate time series data in Python using Pandas. You confirmed that the index was set appropriately, performed queries to subset the data, and practiced identifying and addressing missing values.